




Manifest  
file

YAML or JSON

Describe desired  
state








Desired state/  
record of intent

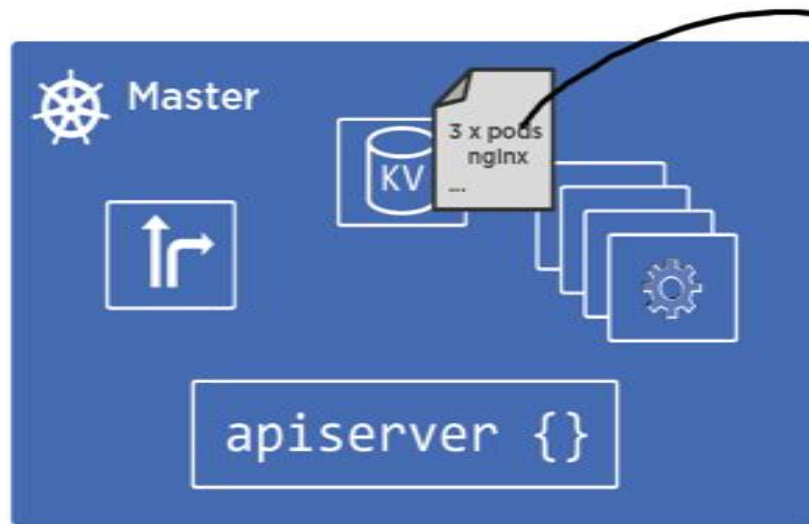
- **3 x nginx pods**



Actual state

- **3 x nginx pods**





Desired state/  
record of intent


- **3 x nginx pods**



Actual state

- **2 x nginx pods**





Desired state/  
record of intent

- **3 x nginx pods**



Actual state

- **3 x nginx pods**



A white line drawing of a pod, which is a small, boat-like structure. It has a flat top and a curved bottom. The word "Pods" is written in white text inside the pod.

Pods



VM



Container



Pod


Atomic units of scheduling





Containers always run  
inside of pods

Pods can have multiple  
containers  
(advanced use-case)




Ring-fenced environment


- Network stack
- Kernel namespaces
- ...

$n$  containers

All containers in pod share the pod environment




# Tight Coupling




# Loose Coupling



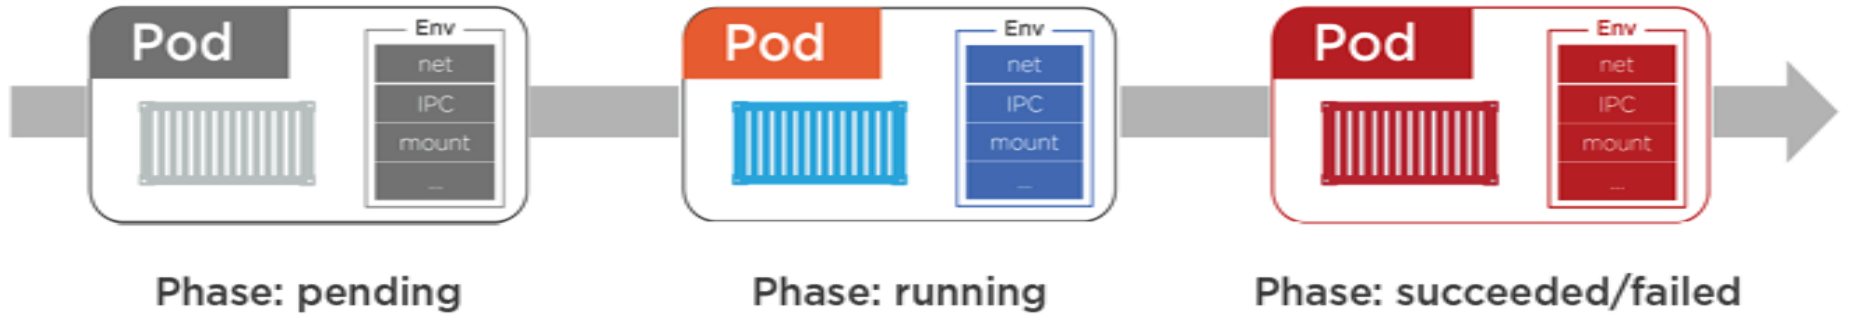
# Pods and Scaling



# Multi-container Pods

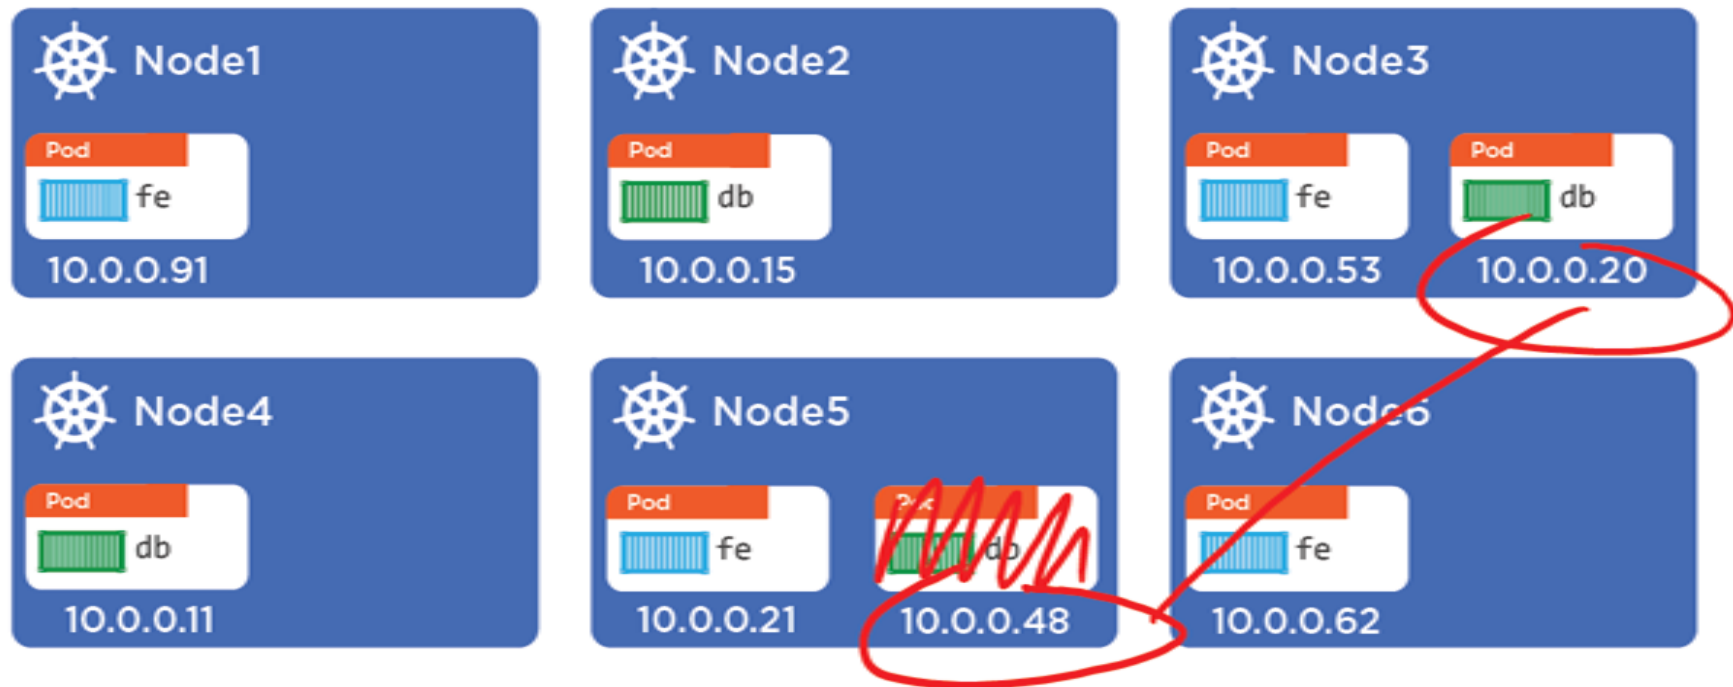


# Pod Lifecycle

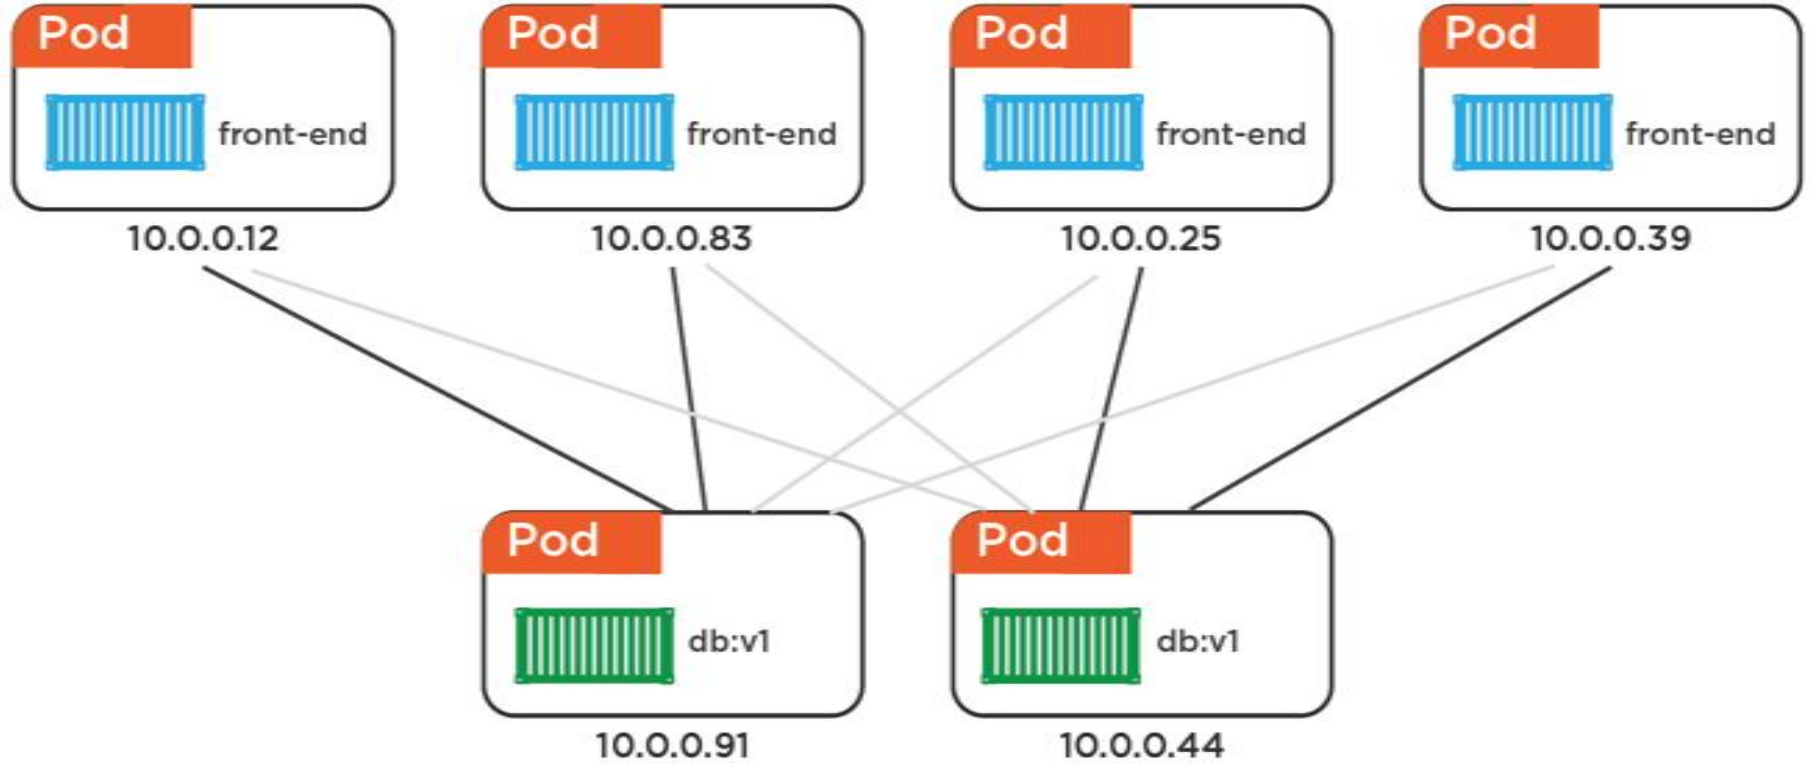


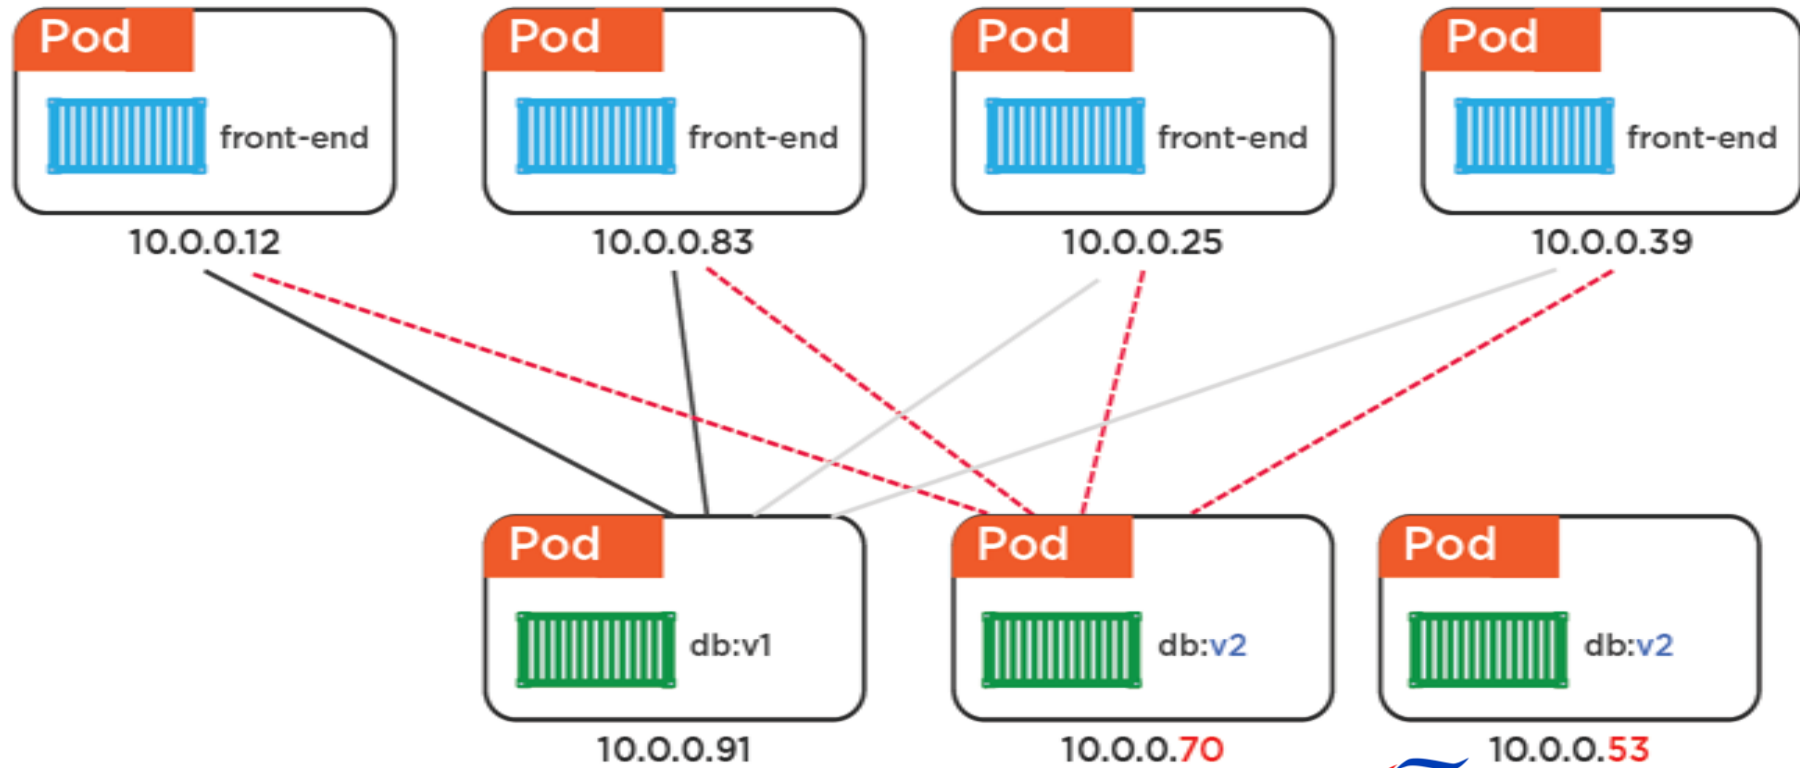
Services

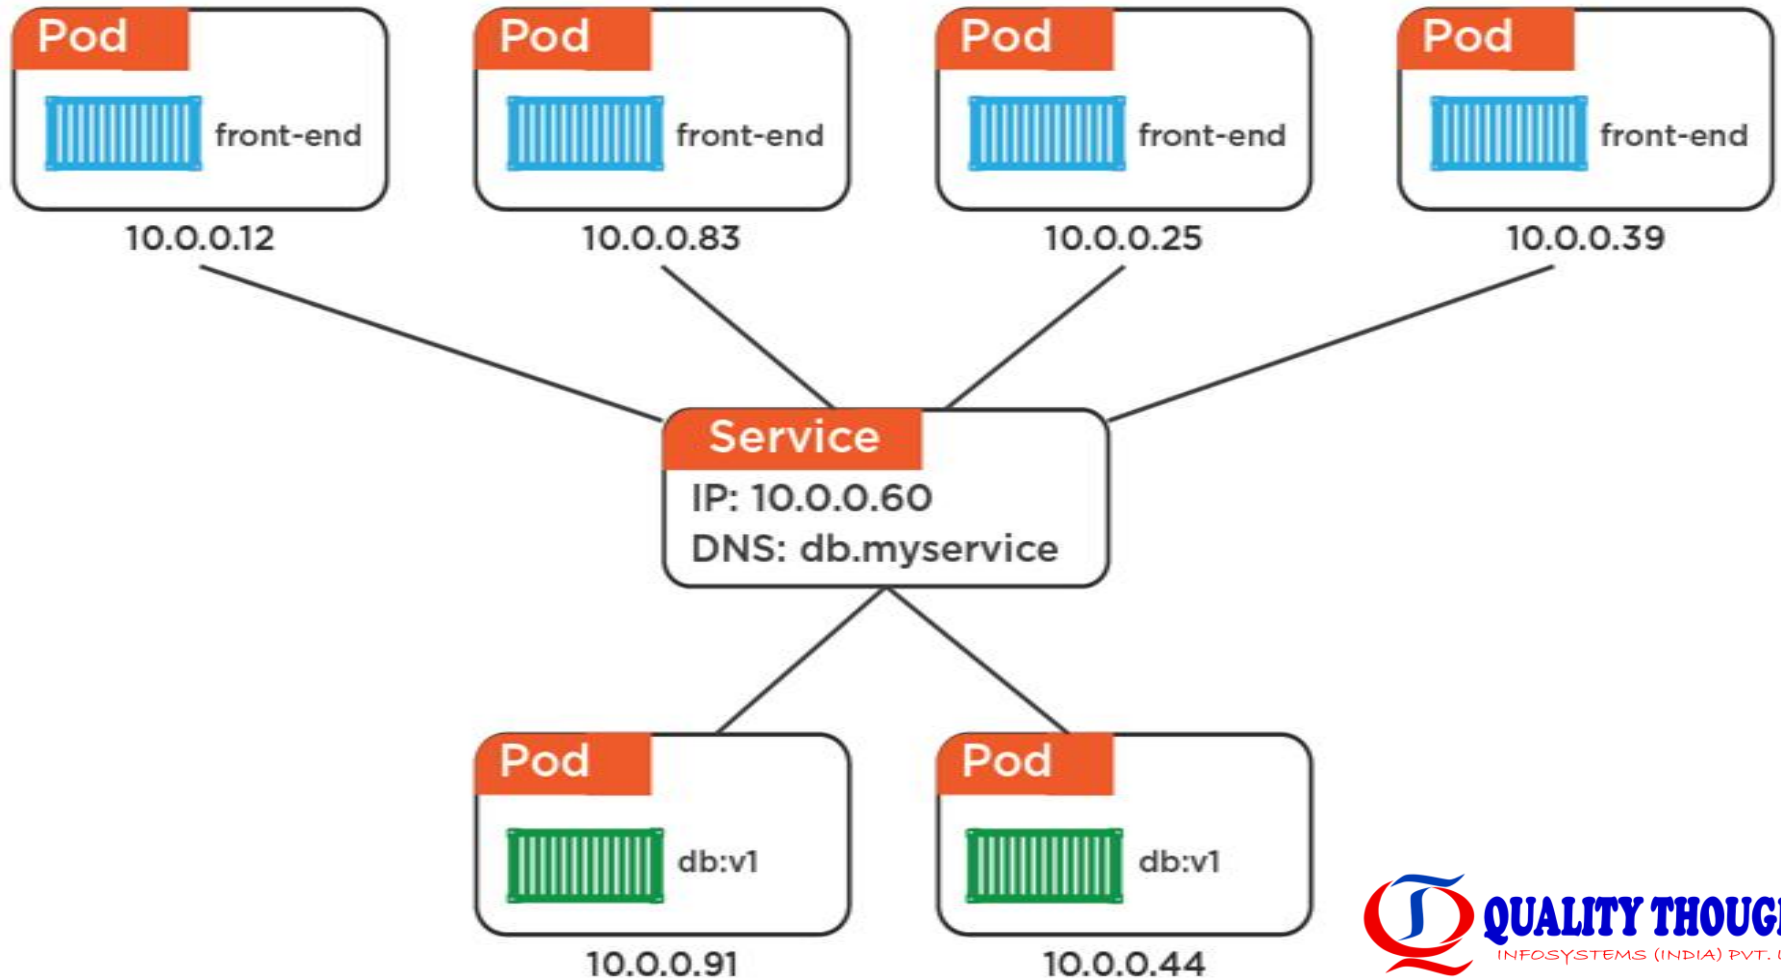








Every new pod gets a new IP = IP churn!






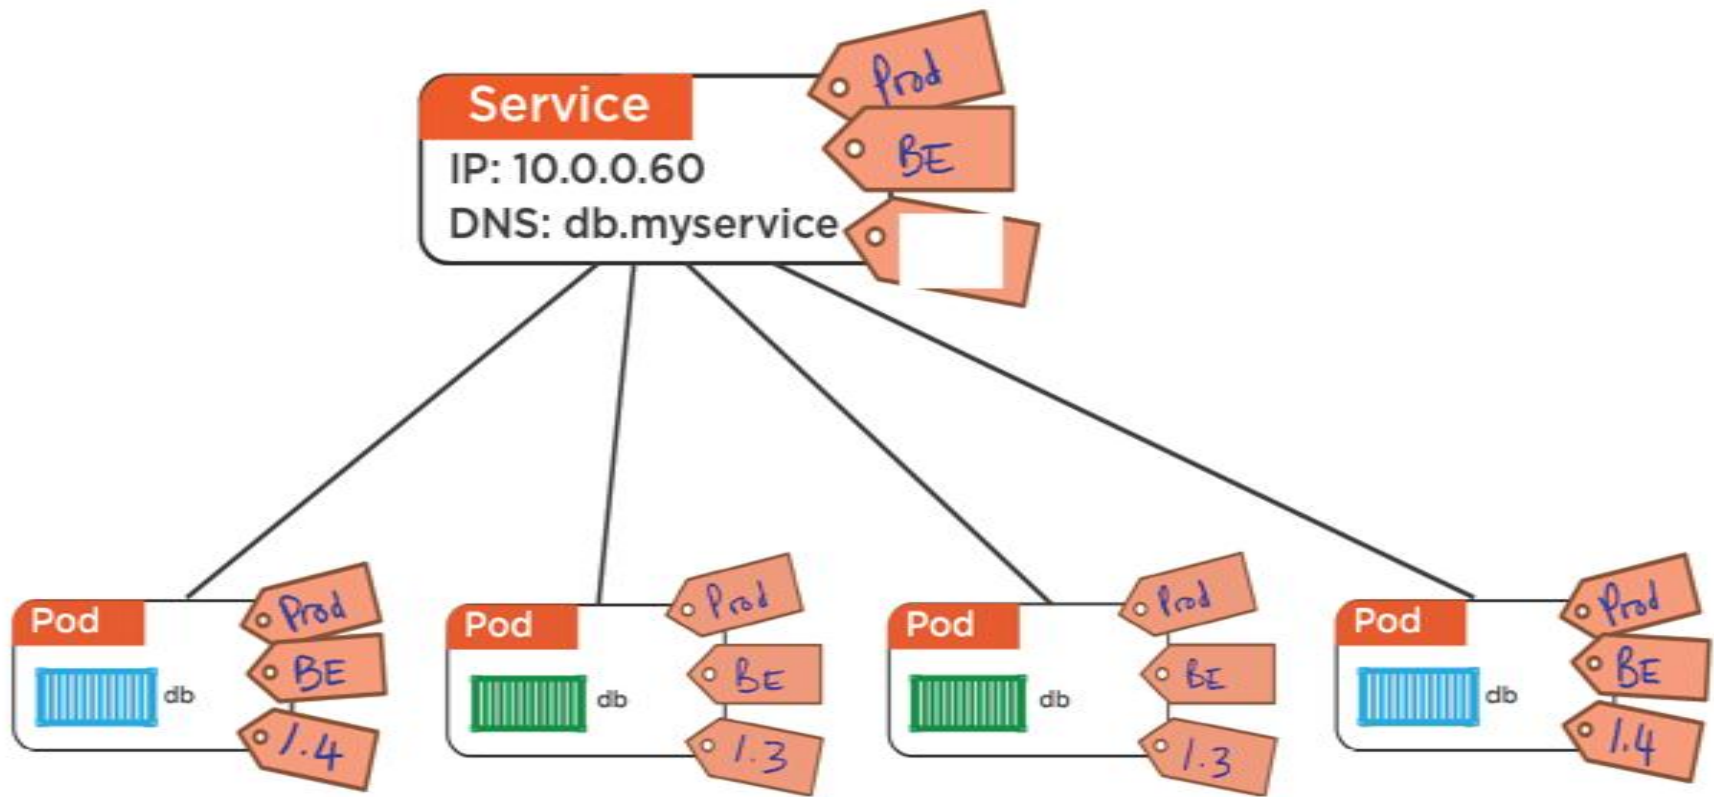




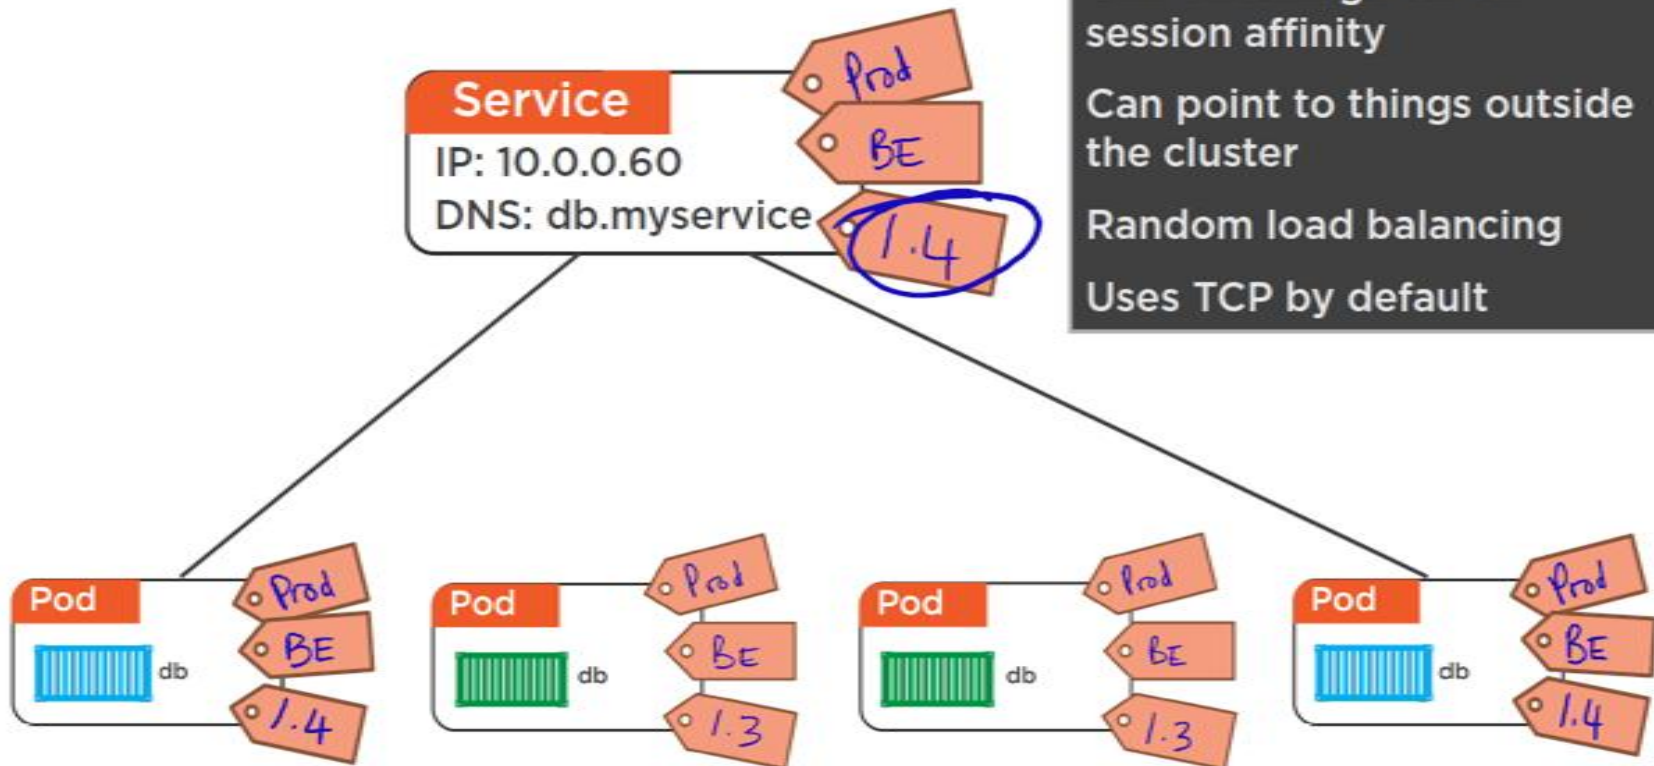












# Deployments

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: xyz
spec:
  replicas: 4
```




Master




apiserver {}




Node



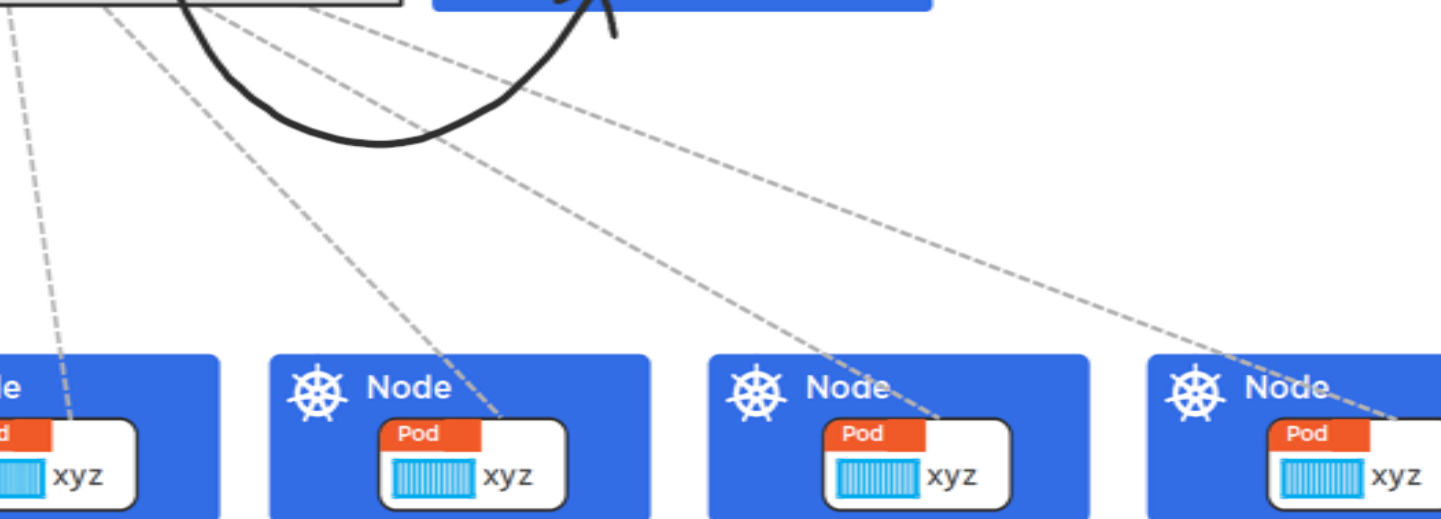

Node



Node



Node



Deployments are all about declarations

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: xyz
spec:
  replicas: 4
```

REST objects

Self documenting

Deployed via YAML or  
JSON manifests

Spec-once deploy-many

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: xyz
spec:
  replicas: 4
```

Simple rolling updates  
and rollbacks

Add features to  
Replication Controllers  
(Replica Sets) *RC v2*

Versioned

Deployed via the  
apiserver

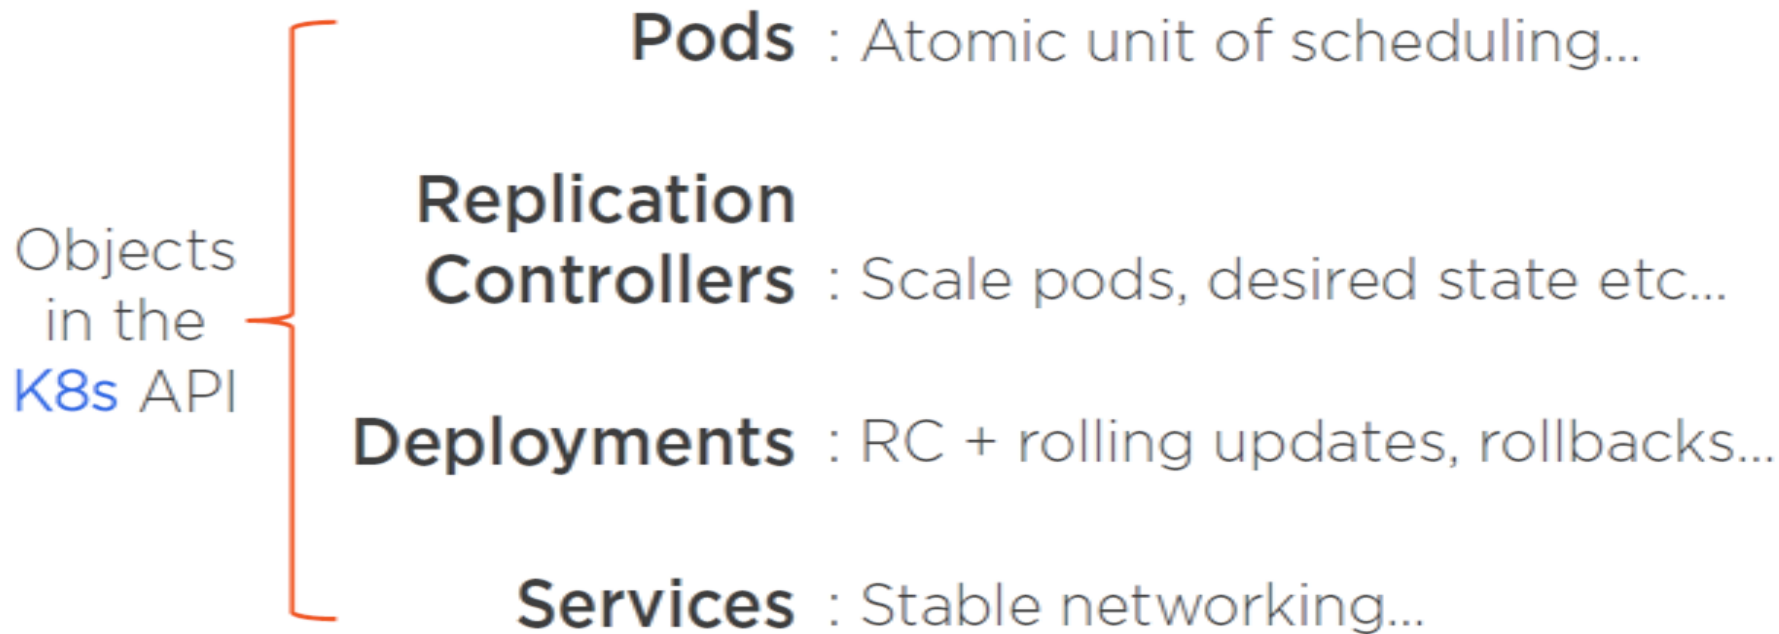
```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: xyz
spec:
  replicas: 4
```

Simple rolling updates  
and rollbacks

Multiple concurrent versions

- Blue-green deployments
- Canary releases

Simple versioned rollbacks



# Installing Kubernetes





VS





VS



Simplest

Most  
involved



Minikube



Google  
Container  
Engine  
(GKE)



AWS  
Provider




Manual  
install

# Minikube








**minikube**



**DFM**  
**DFW**

Best way of spinning up a local environment





`$kubectl`



Node



Master



apiserver {}

Localkube (binary)



Container runtime



# minikube

Best way of spinning up a local k8s environment





# INSTALLING MINIKUBE ON WINDOWS 10

# Get the kubectl.exe binary from this URL and copy it into your %PATH%

<https://storage.googleapis.com/kubernetes-release/release/v1.6.0/bin/windows/amd64/kubectl.exe>

# Get the Minikube installer from GitHub and install

<https://github.com/kubernetes/minikube/releases>

minikube version

minikube start --vm-driver=hyperv --kubernetes-version="v1.6.0"

kubectl get nodes

minikube dashboard

# Manually Installing k8s with `kubeadm`

## Pre-reqs



## Every node (master and minions) needs:

- Docker (or rkt)
- Kubelet
- Kubeadm
- Kubectl
- CNI

# INSTALLING K8S WITH KUBEADM

Install kubeadm as mentioned in

<https://kubernetes.io/docs/setup/independent/install-kubeadm/> or

```
apt-get update && apt-get install -y apt-transport-https
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -
cat <<EOF >/etc/apt/sources.list.d/kubernetes.list
deb http://apt.kubernetes.io/ kubernetes-xenial main
EOF
apt-get update
apt-get install -y docker.io kubeadm kubectl kubelet kubernetes-cni
kubeadm init (as root)
```

```
kubectl get nodes
kubectl get pods --all-namespaces
kubectl apply --filename https://git.io/weave-kube-1.6
```

```
kubectl get nodes
kubectl get pods --all-namespaces
kubectl get nodes
```


# Working with Pods

# Theory

Hypervisor  
Virtualization

VM

Atomic unit of  
scheduling



Container

Atomic unit of  
scheduling



Pod

Atomic unit of  
scheduling